

Assignment 13.a for **STATGR6103**

submitted to Professor Andrew Gelman

Advait Rajagopal

2 December 2016

1 Question 1

Millions of people in rural Bangladesh are exposed to dangerous levels of arsenic in their drinking water which they get from home wells. Several years ago a survey was conducted in small area of Bangladesh to see if people with high arsenic levels would be willing to switch to a neighbor's well. File <http://www.stat.columbia.edu/gelman/bda.course/wells.dat> has the data; all you need here are the variables switch (1 if the respondent said he or she would switch, 0 otherwise) and arsenic (the concentration in the respondent's home well, with anything over 0.5 considered dangerous). Apply the probit model described in exercise 1 above to predict the probability of switching given arsenic level. The goal is inference for the coefficients a and b.

Given that the goal of the survey is inference about the parameters 'a' and 'b', I first examine the data in order to understand the predictor and the dependent variable. Table 1 displays summary statistics of these two variables.

Table 1: Data summary

stat	switch	arsenic
min	0	0.51
mean	0.57	1.66
max	1	9.65

The dependent variable (switch) is a binary response variable, where people indicate their willingness to switch to a neighbor's well. The predictor (arsenic) is the level of arsenic in the drinking water. We are asked to implement a probit model to predict probability of switching given arsenic levels. We are asked to motivate the probit model using a latent variable formulation with $z_i \sim N(\alpha + \beta x_i, 1)$ such that;

$$y_i = \begin{cases} 1, & z_i > 0 \\ 0, & \text{otherwise} \end{cases}$$

The probability that $y_i = 1$ is given by the probit regression below;

$$\begin{aligned} Pr(y_i = 1|x_i) &= \Phi(z_i) \\ &= \Phi(\alpha + \beta x_i) \end{aligned}$$

Where Φ is the cumulative distribution function (CDF) of the standard normal distribution and our goal is to estimate coefficients or parameters α and β using a probit regression, with three techniques. I program an HMC and allow it to estimate the parameters of the probit model, then I run the same model in Stan and finally estimate the parameters using a variational inference algorithm and compare the results.

Since the outcome variable is binary, I use a binomial likelihood function with the following analytical form;

$$p(y|\alpha, \beta, x) \propto \prod_{i=1}^N (\Phi(\alpha + \beta x_i))^{y_i} (1 - \Phi(\alpha + \beta x_i))^{1-y_i}$$

I also consider noninformative priors on α and β such that;

$$p(\alpha, \beta) \propto 1$$

Thus the posterior density is given by;

$$\begin{aligned} p(\alpha, \beta|y, x) &\propto p(\alpha, \beta)p(y|\alpha, \beta, x) \\ &= \prod_{i=1}^N (\Phi(\alpha + \beta x_i))^{y_i} (1 - \Phi(\alpha + \beta x_i))^{1-y_i} \end{aligned}$$

1.1 Part A

Program Hamiltonian Monte Carlo for the probit model described in the previous exercise, again using the latent-variable formulation (so you are jumping in a space of $n+2$ dimensions). Tune the algorithm and run to approximate convergence.

I program an HMC and run 4 chains with 1000 iterations. I set the mass matrix to scale with the variance of the posterior density and tune the ϵ and L parameters in order to get approximately 65% convergence. Table 2 shows the posterior intervals as estimated by the HMC algorithm

Table 2: Posterior Distributions: HMC in R

	mean	se.mean	sd	2.5%	25%	50%	75%	97.5%	n.eff	Rhat
α	-0.2	0	0	-0.3	-0.2	-0.2	-0.2	-0.1	572	1
β	0.2	0	0	0.2	0.2	0.2	0.2	0.3	714	1

1.2 Part B

Check your results by running Stan.

I fit a probit model in Stan and get the results summarized in Table 3.

Table 3: Posterior Distributions: Stan

	mean	se.mean	sd	2.5%	25%	50%	75%	97.5%	n.eff	Rhat
α	-0.18	0.00	0.04	-0.26	-0.21	-0.18	-0.15	-0.11	839	1
β	0.23	0.00	0.02	0.19	0.21	0.23	0.24	0.27	869	1

1.3 Part C

Compare to the variational Bayes inferences for α and β from the previous above.

I program a variational Bayes algorithm to estimate the parameters α and β . The posterior intervals of the parameters are displayed in Table 4. I make some plots to compare the inferences from the

Table 4: Posterior Distributions: Stan

	mean	se.mean	sd	2.5%	25%	50%	75%	97.5%
α	-0.178	0.00	0.03	-0.29	-0.21	-0.18	-0.16	-0.11
β	0.227	0.00	0.03	0.20	0.21	0.23	0.25	0.27

3 approaches, namely HMC, Stan and variational inference. Figure 1 shows a scatterplot of samples drawn from the posterior densities for each of the 3 approaches.

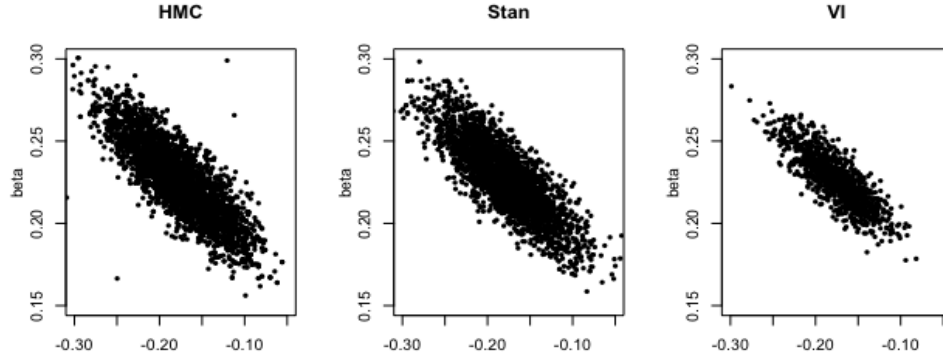


Figure 1: *The results from all the approaches are quite similar.*

Figure 2 and 3 show the marginal posterior density of the parameters α and β as estimated by the 3 approaches, respectively.

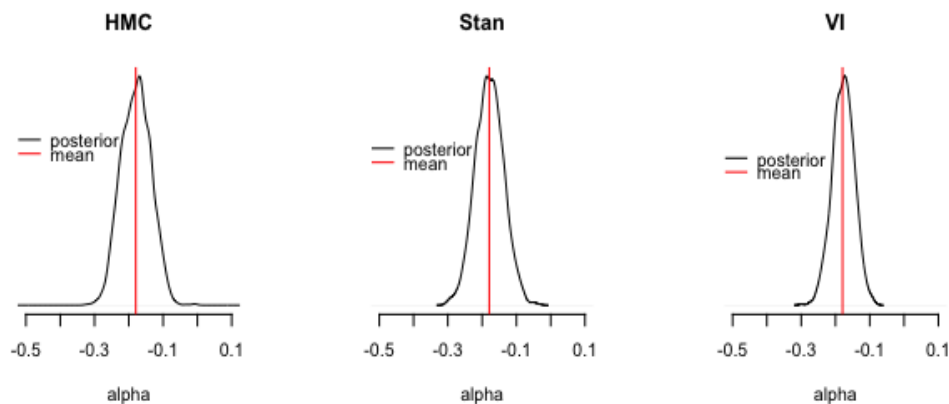


Figure 2: All 3 approaches estimate a similar posterior density for α , the mean of the marginal posterior distribution is marked by the red line.

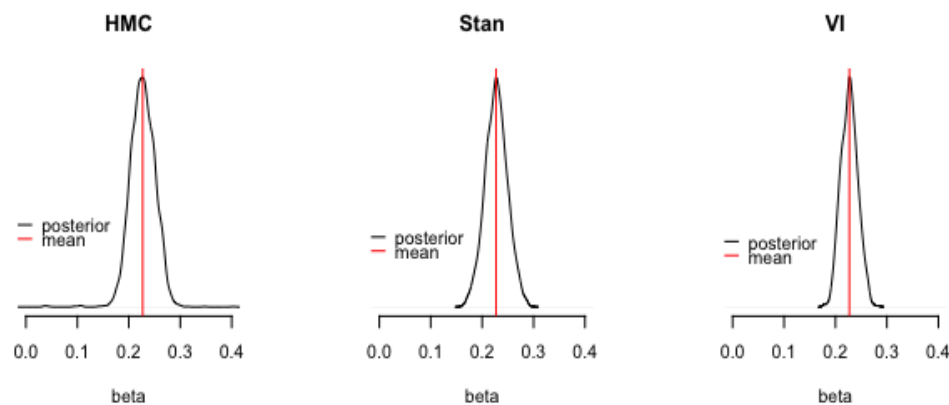


Figure 3: All 3 approaches estimate a similar posterior density for β , the mean of the marginal posterior distribution is marked by the red line.

We see that the results from HMC, Stan and variational Bayes algorithm are very similar and comparable.

2 Code

2.1 R code

Listing 1: R code

```
rm(list = ls())
setwd("/Users/Advait/Desktop/New_School/Fall16/BDA/Class24")
library(rstan)
rstan_options(auto_write = TRUE)
options(mc.cores = parallel::detectCores())
install.packages("data.table", dependencies = T)
library(data.table)
library ( plyr )
library ( arm )
library ( msm )
library ( mvtnorm )
library(LearnBayes)
#####
df <- fread("http://www.stat.columbia.edu/
~gelman/bda.course/wells.dat")
str(df)
names(df) <- c("s.no", "switch", "arsenic", "dist", "assoc", "educ")
summary(df)
###
df_final <- data.frame(df$switch, df$arsenic)
class(df_final)
str(df_final)
##
#Stan
y <- df_final$df.switch
x <- df_final$df.arsenic
N <- length(df_final$df.switch)
N
stanc("13a.stan")
fit_stan <- stan("13a.stan", data = list("N", "y", "x"),
               iter = 2000, chains = 3)
print(fit_stan)

# plot(ext_stan$y_rep)
# points(ext_stan$y_rep, col = "red")
# plot(colMeans(ext_stan$y_rep), y, pch = 16, cex = 0.7)
# abline(0,1)
```

Listing 2: R code contd.

```
##
##PartA
#HMC
#log posterior function
n <- 1
log_post <- function(param, x , y){
  param_prior <- dunif(param,-100,100, log = T)
  log_prior <- sum(param_prior)
  log_likelihood <- sum(dbinom(y, n,
  prob = pnorm(param[1] + param[2]*x),log = T))
  return(log_prior + log_likelihood)
}
log_post(c(0.1,-0.3),x,y)

#Numerical gradient
gradient_num <- function(param, x, y){
  d <- length(param)
  e <- 0.0001
  diff <- rep(NA,d)
  for(k in 1:d){
    th_hi <- param
    th_lo <- param
    th_hi[k] <- param[k] + e
    th_lo[k] <- param[k] - e
    diff[k] <- (log_post(th_hi,x,y) - log_post(th_lo,x,y))/(2*e)
  }
  return(diff)
}
gradient_num(c(-0.001,-.04),x,y)
###
#HMC iter
hmc_iteration <- function(param,x,y, epsilon,L,M){
  M_inv <- 1/M
  d <- length(param)
  # Sample 10 points randomly from a normal
  #distribution with mean = 0 and standard deviation = sqrt(M)
  phi <- rnorm(d , 0, sqrt(M))
  param_old <- param
  log_p_old <- log_post(param,x,y) - 0.5*sum(M_inv*phi^2)
  phi <- phi + 0.5*epsilon*gradient_num(c(param), x, y)
```

Listing 3: R code contd.

```

for (l in 1:L){
  param <- param + epsilon*M_inv*phi
  phi <- phi + (if (l==L)0.5 else 1)*epsilon*
  gradient_num(c(param),x,y)
}
phi <- -phi
log_p_star <- log_post(c(param),x,y) - 0.5*sum(M_inv*phi^2)
r <- exp(log_p_star - log_p_old)
if(is.nan(r)) r <- 0
p_jump <- min(r,1)
param_new <- if(runif(1) < p_jump) param else param_old
return (list (param = param_new, p_jump = p_jump))
}
##
##HMC run
hmc_run <- function (starting_values, iter, epsilon_0, L_0, M) {
  # Get the number of rows and store in chains
  chains <- nrow (starting_values)
  # The number of parameters that you have in the starting values
  d <- ncol (starting_values)
  # Create space to store iterations of the parameters
  sims <- array (NA, c(iter, chains, d),
                 dimnames=list (NULL, NULL,
                                colnames (starting_values)))
  warmup <- 0.5*iter
  p_jump <- array (NA, c(iter, chains))
  for (j in 1:chains){
    param <- starting_values[j,]
    for (t in 1:iter){
      epsilon <- runif (1, 0, 2*epsilon_0)
      L <- ceiling (2*L_0*runif(1))
      temp <- hmc_iteration (param,x,y, epsilon,L,M)
      p_jump[t,j] <- temp$p_jump
      sims[t,j,] <- temp$param
      param <- temp$param
    } }
  monitor (sims, warmup)
  cat ("Avg_acceptance_probs:",
       fround(colMeans(p_jump[(warmup+1):iter,]),2),"\n")
  return (list (sims=sims, p_jump=p_jump))
}

```

Listing 4: R code contd.

```
##RUN it
parameter_names <- c (paste ("beta[",1:2,"]",sep=""))
d <- 2
chains <- 4

#
log_post_interim <- function(param){
  log_prior <- 0
  log_likelihood <- sum(dbinom(y, n,
    invlogit(param[1] + param[2]*x),log = T))
  return(log_prior + log_likelihood)
}
lat <- laplace(log_post_interim, c(0.5,-0.5))
M <- ginv(lat$var)
mass_vector <- c(0.005, 0.001)

#Starts
starts <- array (NA,c(chains,d),
  dimnames=list(NULL,parameter_names))
for (j in 1:chains){
  starts[j,1] <- rnorm (1,0,0.5)
  starts[j,2] <- rnorm (1,0,0.5)
}
fit_hmc1 <- hmc_run(starting_values = starts,
  iter = 1000,
  epsilon_0 = 0.0005,
  L_0 = 10,
  M = mass_vector)
```


Listing 5: R code contd.

```
#####
##Variational Bayes

N <- 3020
X <- matrix(c(rep(1, N), x), nrow = N, ncol = 2)

#Lower bound function
lower.bound2 <- function(param, X, stars, sigs, draws){
  ab <- t(X)\%*\%X\%*\%(param\%*\%t(param) + sigs)
  part1 <- sum(diag(ab))/2
  part2 <- t(param)\%*\%param + sum(diag(sigs))
  part2 <- part2*(beta.mat.prior[1,1])
  part2 <- part2/2
  + (1/2)*log(det(solve(beta.mat.prior)))
  + (length(param)/2)*log(2*pi)
  part3 <- t(stars)\%*\%stars
  part4 <- length(param)/2
  + (1/2)*log(det(sigs))
  + (length(param)/2)*log(2*pi)
  bounds <- part1 + part2 + part3/2 + part4
  parts <- c(-part1, -part2, part3/2, part4)
  bounds<- list(bounds, parts)
  names(bounds)<- c("bounds", "parts")
  return(bounds)
}

##we will use the lower-bound to monitor convergence,
##stopping the model when the lower-bound drops below 1e-8
func.reg<- function(X, draws){
  beta.mat.prior <- diag(1/10, ncol(X))
  beta.VA <- matrix(NA, nrow=1000, ncol=ncol(X))
  ystars <- rep(0, nrow(X))
  for(j in 1:nrow(X)){
    ystars[j] <- ifelse(draws[j]==1,
                        rtnorm(1, mean=0.5, sd=1, lower=0, upper=Inf),
                        rtnorm(1, mean=-0.5, sd=1, lower=-Inf, upper=0))
  }
  ##we will store the progress of the lower bound on the model
  bounds<- c()
  zz <- 0
}
```

Listing 6: R code contd.

```

##this stores the parts of the lower bound
parts<- matrix(NA, nrow=1000, ncol=4)
j <- 0
##creating a while loop
while(zz == 0){
  j <- j + 1
  ##updating the beta parameters
  beta.VA[j,] <- solve(t(X)\%*\%X + beta.mat.prior)
  \%*\%t(X)\%*\%ystars
  sigs <- solve(t(X)\%*\%X + beta.mat.prior)
  ##computing the inner product of the
  ##covariates current estimates of the coefficients
  stars <- X\%*\%beta.VA[j,]
  denom1 <- pnorm(-stars)
  num1 <- dnorm(-stars)
  ##now, computing the expected value for each
  ##individual's augmented data, given
  ####current estimates of the approximating
  ####distribution on the coefficients
  ####ystars[which(draws==0)]<-stars[draws==0]+
  #####-num1[which(draws==0)]/denom1[which(draws==0)]
  ####ystars[which(draws==1)]<-stars[draws==1]+
  #####num1[which(draws==1)]/(1-denom1[which(draws==1)])
  ####calculating the lower bound
  ####trial<-lower.bound2(beta.VA[j,],X,ystars,sigs,draws)
  ####bounds[j]<-trial$bounds
  ####parts[j,]<-trial$parts
  ####if(j>1){
  #####monitor convergence
  #####ab<-abs(bounds[j]-bounds[j-1])
  #####if(ab<1e-8){
  #####zz<-1}
  ####}
  ##}
  ##the information to be returned, after convergence
  stuff<-list(bounds,beta.VA[j,],sigs)
  names(stuff)<-c("bound","param","sigma")
  return(stuff)
}

```

Listing 7: R code contd.

```

example.run<- func.reg(X, y)
example.run
betas <- rmvnorm(10^3,
                mean = example.run$param,
                sigma = example.run$sigma)
#####

par(mfrow = c(1, 3),
    mgp = c(2,1,0),
    mar = c(2,3,4,2))
plot(fit_hmc1$sims[, ,1], fit_hmc1$sims[, ,2], ylim = c(0.15, 0.30),
     xlim = c(-0.3, -0.05), pch = 16, cex = 0.7,
     main = "HMC", xlab = "alpha", ylab = "beta")
plot(ext_stan$alpha, ext_stan$beta, pch = 16, cex = 0.7,
     main = "Stan", xlab = "alpha", ylab = "beta", ylim = c(0.15, 0.30),
     xlim = c(-0.3, -0.05))
plot(betas[,1], betas[,2], pch = 16, cex = 0.7,
     main = "VI", xlab = "alpha", ylab = "beta", ylim = c(0.15, 0.30),
     xlim = c(-0.3, -0.05))
#####

par(mfrow = c(1, 3))
plot(density(fit_hmc1$sims[, ,1]), main = "HMC",
     xlab = "alpha", ylab = NA,
     bty = "n", yaxt = "n", xlim = c(-0.5, 0.1))
abline(v = mean(fit_hmc1$sims[, ,1]), col = "red")
legend(-0.55, 7, legend = c("posterior", "mean"), lty = c(1, 1),
     col = c("black", "red"), bty = "n")

##
plot(density(ext_stan$alpha), main = "Stan",
     xlab = "alpha", ylab = NA,
     bty = "n", yaxt = "n", xlim = c(-0.5, 0.1))
abline(v = mean(fit_hmc1$sims[, ,1]), col = "red")
legend(-0.55, 7, legend = c("posterior", "mean"), lty = c(1, 1),
     col = c("black", "red"), bty = "n")

#
plot(density(betas[,1]), main = "VI",
     xlab = "alpha", ylab = NA,
     bty = "n", yaxt = "n", xlim = c(-0.5, 0.1))
abline(v = mean(fit_hmc1$sims[, ,1]), col = "red")
legend(-0.55, 8.5, legend = c("posterior", "mean"), lty = c(1, 1),
     col = c("black", "red"), bty = "n")

```

Listing 8: R code contd.

```
#####
par(mfrow = c(1, 3))
plot(density(fit_hmc1$sims[,2]),main = "HMC",
      xlab = "beta", ylab = NA,
      bty = "n", yaxt = "n", xlim = c(0, 0.4))
abline(v = mean(fit_hmc1$sims[,2]), col = "red")
legend(-0.05,7, legend = c("posterior", "mean"),lty = c(1,1),
      col = c("black", "red"), bty = "n")
##
plot(density(ext_stan$beta),main = "Stan",
      xlab = "beta", ylab = NA,
      bty = "n", yaxt = "n", xlim = c(0, 0.4))
abline(v = mean(fit_hmc1$sims[,2]), col = "red")
legend(-0.05,7, legend = c("posterior", "mean"),lty = c(1,1),
      col = c("black", "red"), bty = "n")
#
plot(density(betas[,2]),main = "VI",
      xlab = "beta", ylab = NA,
      bty = "n", yaxt = "n", xlim = c(0, 0.4) )
abline(v = mean(fit_hmc1$sims[,2]), col = "red")
legend(-0.05,8.5, legend = c("posterior", "mean"),lty = c(1,1),
      col = c("black", "red"), bty = "n")
```

2.2 Stan Code

Listing 9: Stan Code

```
data{
  int N;
  int y[N];
  real x[N];
}
parameters{
  real alpha;
  real beta;
}
model{
  for(i in 1:N)
    y[i] ~ bernoulli(Phi(alpha + beta * x[i]));
// }
// generated quantities{
//   real y_rep[N];
//   for(i in 1:N)
//     y_rep[i] = bernoulli_rng(Phi(alpha + beta * x[i]));
// }
```